SOP-SP-002: New Service Principal Provisioning (Azure Entra ID + Databricks)

Owner: MLOps Platform Team | Version 1.0 | Sample SOP for prototype

1. Purpose and Scope

- This SOP covers creating a brand-new service principal (application identity) in Azure Entra ID and registering it in one or more Databricks workspaces.
- Service principals are used by automated jobs, workflows, pipelines and CI/CD instead of personal accounts.

2. When to use this SOP

- A team needs a new automated identity for a scheduled job, batch inference, Databricks Workflow or Jenkins pipeline.
- A production job currently runs with personal user credentials and must move to a service principal.

3. Not covered by this SOP

- Adding people or groups to an EXISTING service principal: use SOP-SP-003.

4. Required information in the ticket

- Proposed name following convention sp--.
- Target environments / workspaces.
- Named owner (a permanent employee) and cost center.
- Purpose of the identity and whether a client secret is needed.

5. Validation rules

- Name must match sp-- and be unique in the tenant.
- Owner must be a permanent employee, not a contractor.
- Client secrets have a 180-day expiry and are stored only in Azure Key Vault.

6. Approval

- Job Owner (JO) of the consuming job approves the creation.
- Data Product Owner (DPO) approval is needed only if data grants are requested in the same ticket.

7. Implementation

- Repository: grey-infra.
- Branch: feature/-new-sp created from master.
- File: terraform/service_principals/.tf using module sp_with_databricks_registration.

8. Deployment

- Jenkins job: grey-infra-terraform-plan on PR, grey-infra-terraform-apply after merge to master.

9. Post-deployment verification

- Confirm the application exists in Entra ID and the service principal is listed in each Databricks workspace.
- Confirm the secret is in Key Vault and share the Key Vault reference (never the secret) on the ASPEN ticket.